GRU Λ < MWPM Λ (5→7: 1.81 vs 1.76; 7→9: 0.56 vs 1.89); GRU Λ < 1 from d7 to d9

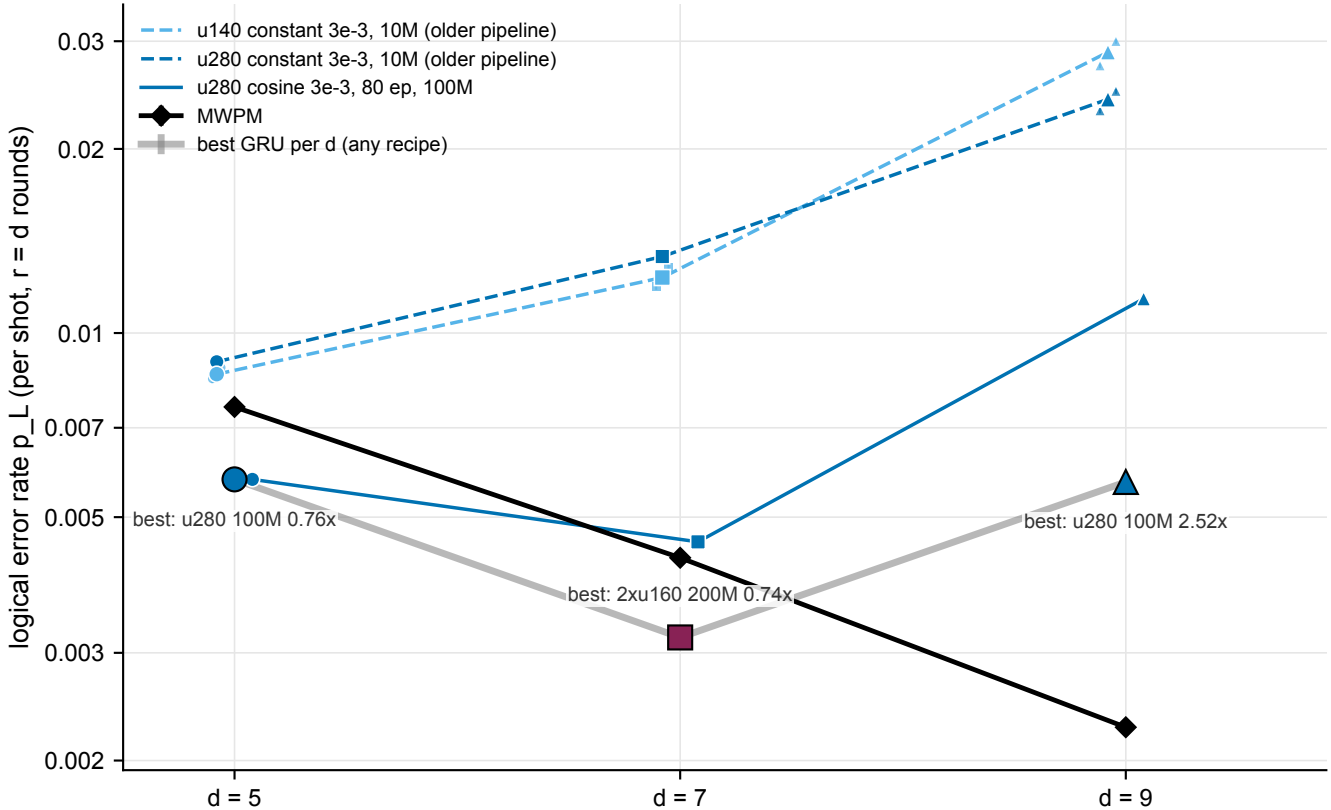

series	Lambda 5->7 [95% CI]	Lambda 7->9 [95% CI]
MWPM	1.76 [1.72, 1.81]	1.89 [1.82, 1.96]
best GRU per d	1.81 [1.76, 1.87]	0.56 [0.54, 0.57]
u140 constant 3e-3, 10M (older pipeline)	0.69 [0.67, 0.72]	0.43 [0.41, 0.45]
u280 constant 3e-3, 10M (older pipeline)	0.67 [0.66, 0.69]	0.55 [0.54, 0.58]
u280 cosine 3e-3, 80 ep, 100M	1.27 [1.23, 1.30]	0.40 [0.39, 0.41]

Lambda = p_L(d) / p_L(d+2), bootstrap 95% CI. Grey = best GRU at each d (any recipe within the u280 parameter budget); coloured lines use one recipe at every d. Eval block: 1.8M shots per d, same MWPM for every run at a given d. Bars: Wilson 95%. Data: GRU_all_runs_2026-10-01.csv.