Unseen p_L is 10-36% above own-training-shot p_L; 300 epochs on 2M shots memorizes

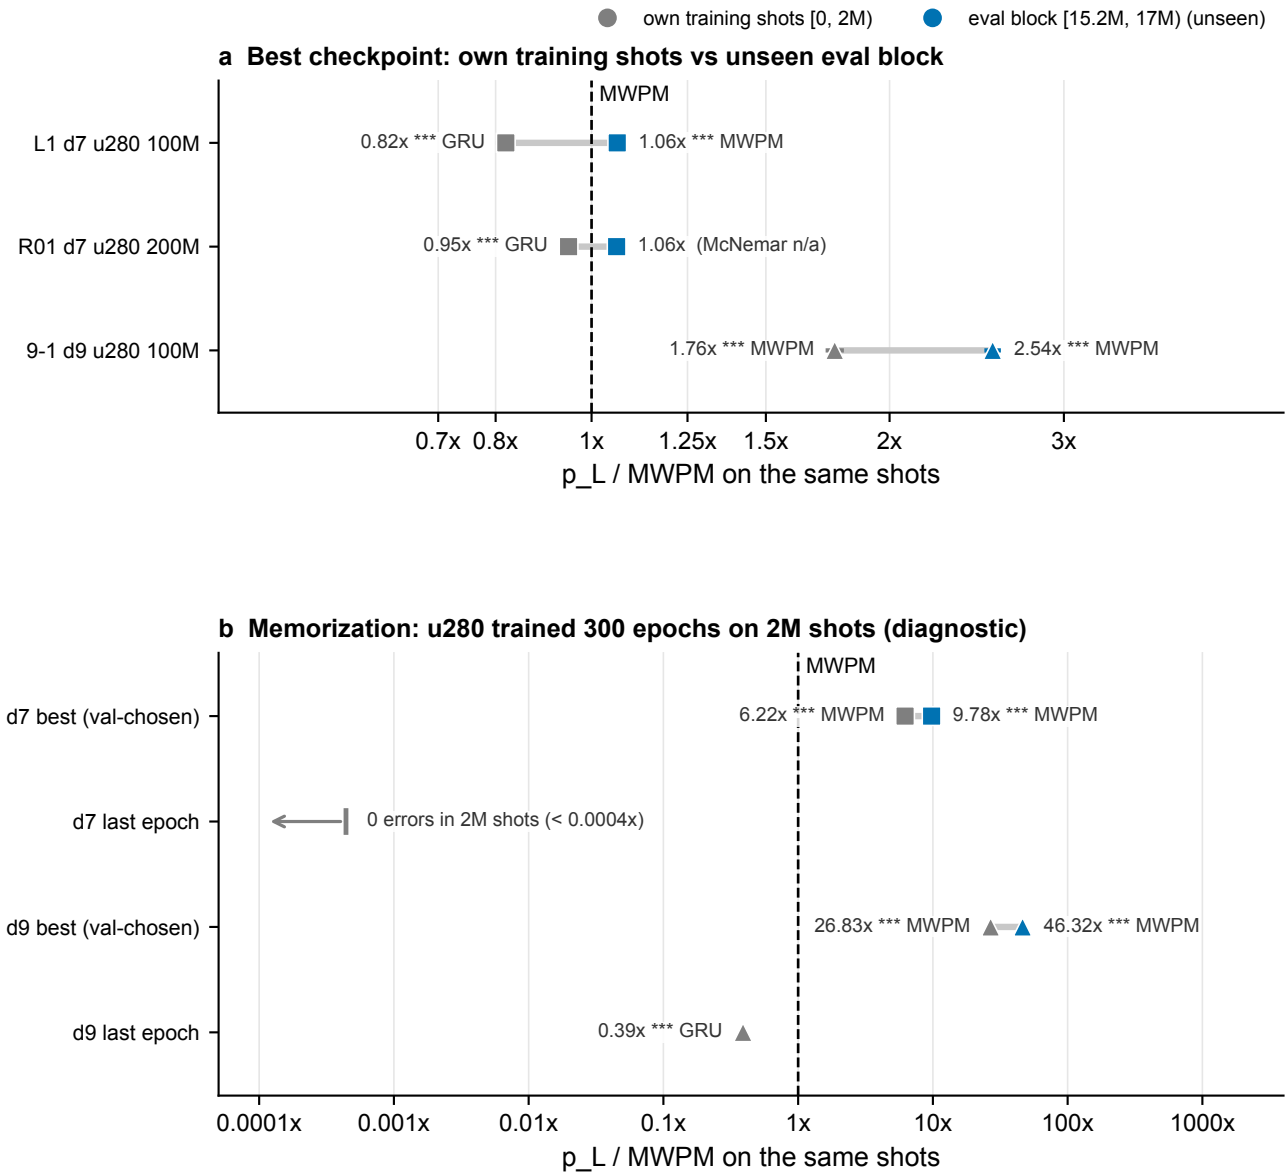

Last-epoch weights on unseen validation shots: d7 p_L 0.051; d9 p_L 0.132 (vs. 0.004-0.005 for the 100M runs)

(a) Best checkpoint on its own training shots vs the eval block; stars = McNemar vs MWPM (*** p<1e-3).
(b) Diagnostic: trained 300 ep on 2M shots only. Eval block: 1.8M shots per d, same MWPM for every run at a given d. Bars: Wilson 95%. Data: GRU_all_runs_2026-10-01.csv, figures/data/extra_scores.csv.